

## Lab2: send and receive packets with DPDK

521021910454 黄培正

### 1. 问题解答

- Q1: What's the purpose of using hugepage?
  - 减少 TLB 缓存项的使用，从而大大降低 TLB Miss 的概率。
  - 减少页表的级数，从而提升查询页表的效率。
- Q2: Take examples/helloworld as an example, describe the execution flow of DPDK programs?

答：helloworld 代码如下所示。

```
int main(int argc, char **argv)
{
    int ret;
    unsigned lcore_id;

    ret = rte_eal_init(argc, argv);
    if (ret < 0)
        rte_panic("Cannot init EAL\n");

    /* call lcore_hello() on every worker lcore */
    RTE_LCORE_FOREACH_WORKER(lcore_id) {
        rte_eal_remote_launch(lcore_hello, NULL, lcore_id);
    }

    /* call it on main lcore too */
    lcore_hello(NULL);
    rte_eal_mp_wait_lcore();

    /* clean up the EAL */
    rte_eal_cleanup();
    return 0;
}
```

控制流如下：

1. 调用 `rte_eal_init`，初始化基础运行环境，若初始化失败则报错。
  2. 对多核运行初始化。即遍历 EAL 指定可以使用的 lcore，然后通过 `rte_eal_remote_launch` 在每个 lcore 上，启动被指定的线程。
  3. 运行当前线程的函数 `lcore_hello`。
  4. 主线程等待从线程结束执行。
  5. 执行 `rte_eal_cleanup`，释放资源，防止 hugepage 内存泄漏。
- Q3: Read the codes of examples/skeleton, describe DPDK APIs related to sending and receiving packets.

答：以下分别为收包、发包所对应的 API。通过指定的端口与队列，收/发缓存区中的数据。

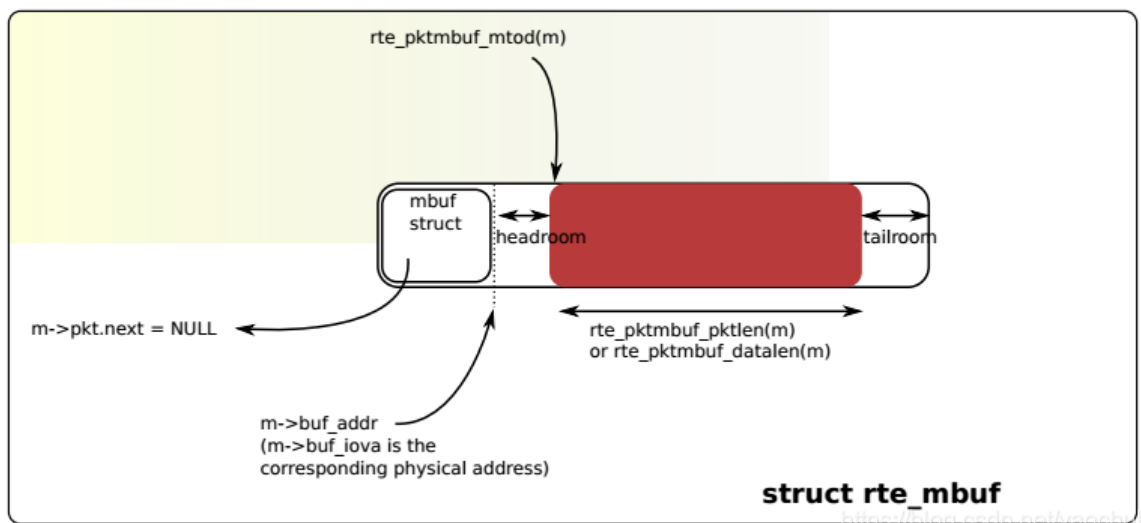
其中，最后一个参数`nb\_pkts`为指定一次函数调用来处理的包的个数。当设置为1时，每次收/发一个包。

dpdk在样例程序，例如skeleton中，使用了burst模式，即收/发包数量为32个。这样可以减少内存访问，提高性能。

```
static inline uint16_t rte_eth_rx_burst(uint8_t port_id, uint16_t queue_id,
    struct rte_mbuf **rx_pkts, const uint16_t nb_pkts)
```

```
static inline uint16_t rte_eth_tx_burst(uint8_t port_id, uint16_t queue_id,
    struct rte_mbuf **tx_pkts, uint16_t nb_pkts)
```

- Q4: Describe the data structure of 'rte\_mbuf'.
- 答：rte\_mbuf的结构如下所示。



- `headroom`为mbuf头部与实际包数据的一段空间，存储控制信息、帧内容、事件等。`headroom`的起始地址保存在`buf_addr`指针中。
- 在`headroom`后为实际数据所占空间。数据帧的长度可通过调用`pkt_len(m)`或`data_len(m)`获得。
- 实际数据后剩余的空间为`tailroom`。通过`headroom`与`tailroom`，可方便应用解封报文。
- `pkt`的`next`字段指向下一个 segment 的地址；`buf_addr`指向`headroom`的起始地址；`rte_pktmbuf_mtod(m)`指向实际 data 的起始地址。
- 此外，还记录了所属的 mempool，时间戳，端口，私有数据大小等信息。

## 2. 检验正确性

主要代码如下：

```

uint16_t port;

/*
 * Check that the port is on the same NUMA node as the polling thread
 * for best performance.
 */
RTE_ETH_FOREACH_DEV(port)
if (rte_eth_dev_socket_id(port) >= 0 &&
    rte_eth_dev_socket_id(port) !=
    (int) rte_socket_id())
    printf("WARNING, port %u is on remote NUMA node to "
           "polling thread.\n\tPerformance will "
           "not be optimal.\n", port);

printf("\nCore %u forwarding packets. [Ctrl+C to quit]\n", rte_lcore_id());

char *package_data = "hello from virtual machine by hpz";
int data_len = strlen(package_data) + 1;

/* Run until the application is quit or killed. */

port = 0; // only use port 0

struct rte_mbuf *buf[1];
buf[0] = rte_pktmbuf_alloc(mbuf_pool); //分配空间
int hdr_len = sizeof(
    struct rte_ether_hdr) + sizeof(
    struct rte_ipv4_hdr) + sizeof(
    struct rte_udp_hdr);
rte_pktmbuf_append(buf[0], data_len + hdr_len);

puts("checkpoint 1");
printf("%d %d\n", data_len, hdr_len);

/* 构造包 ether+ip+udp*/
struct rte_ether_hdr *eth_hdr = rte_pktmbuf_mtod(buf[0],
struct rte_ether_hdr *);
struct rte_ipv4_hdr *ip_hdr = rte_pktmbuf_mtod_offset(buf[0],
struct rte_ipv4_hdr *, sizeof(struct rte_ether_hdr));
struct rte_udp_hdr *udp_hdr = rte_pktmbuf_mtod_offset(buf[0],
struct rte_udp_hdr *, sizeof(struct rte_ether_hdr) + sizeof(struct
rte_ipv4_hdr));

puts("checkpoint 2");

/* 构造ethernet header */
struct rte_ether_addr mac_addr;
rte_eth_macaddr_get(0, &mac_addr);
eth_hdr->ether_type = rte_cpu_to_be_16(RTE_ETHER_TYPE_IPV4);
eth_hdr->s_addr = mac_addr;
eth_hdr->d_addr = mac_addr;

```

```
puts("checkpoint 3");

/* 构造ip header */
ip_hdr->version_ihl = RTE_IPV4_VHL_DEF;
ip_hdr->type_of_service = 0; // unused
ip_hdr->total_length = 62 << 8; // calc by data length
ip_hdr->packet_id = 0;
ip_hdr->fragment_offset = 0;
ip_hdr->time_to_live = 255;
ip_hdr->next_proto_id = 17; // udp
ip_hdr->src_addr = IPv4(58, 247, 22, 200);
ip_hdr->dst_addr = IPv4(202, 120, 40, 85);
ip_hdr->hdr_checksum = rte_ipv4_cksum(ip_hdr);

puts("checkpoint 4");

/* 构造udp header */
udp_hdr->src_port = 8085;
udp_hdr->dst_port = 8080;
udp_hdr->dgram_len = (data_len + sizeof(struct rte_udp_hdr)) << 8;
udp_hdr->dgram_cksum = rte_ipv4_udptcp_cksum(ip_hdr, udp_hdr);

puts("checkpoint 5");

void *tmp = rte_pktmbuf_mtod_offset(buf[0],
void *, hdr_len);
memcpy(tmp, package_data, data_len);

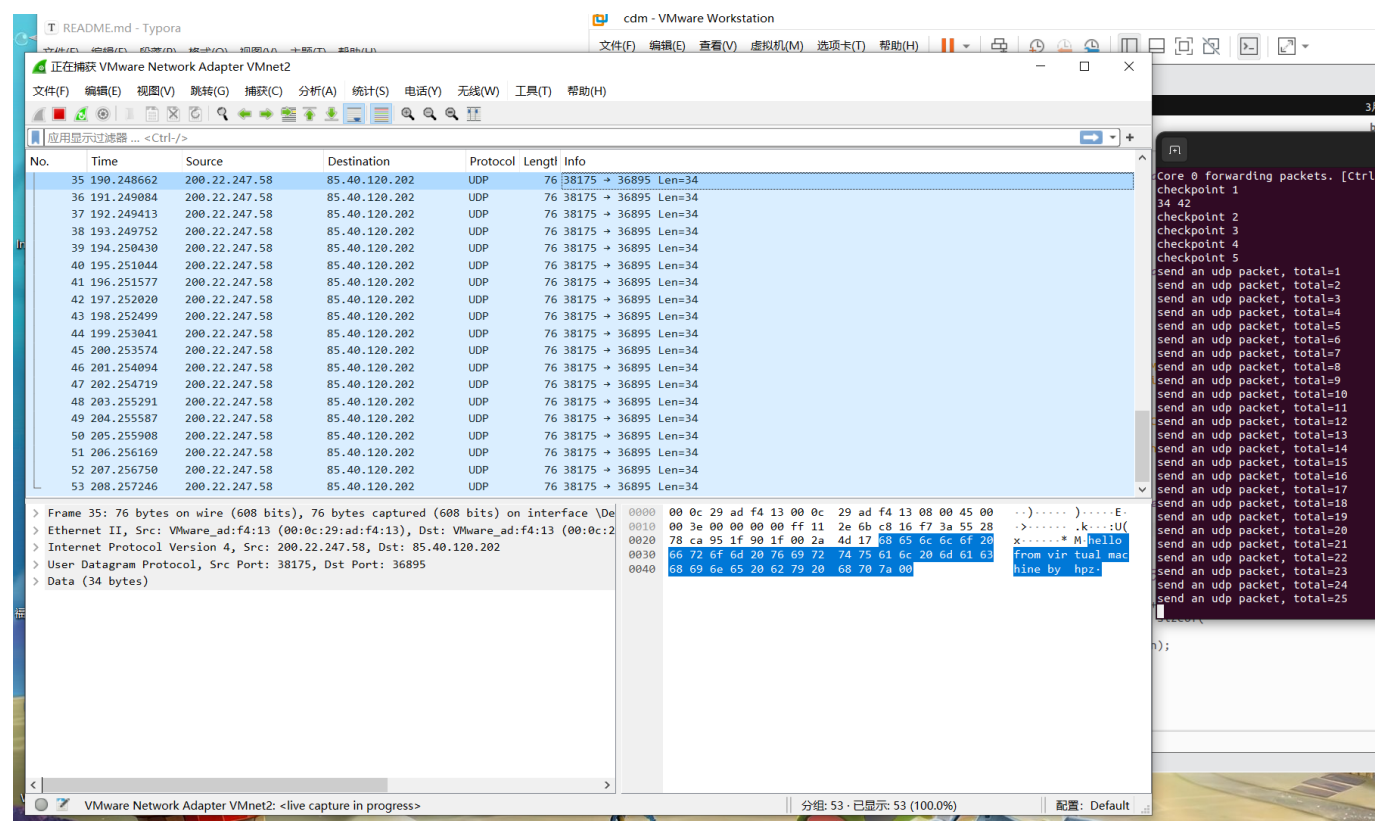
int total = 0;
int interval = 1;
for (;;) {
    const uint16_t nb_rx = rte_eth_tx_burst(port, 0, buf, 1); //每次发一个包
    total++;
    printf("send an udp packet, total=%d\n", total);
    sleep(interval);
}

/* Free packets. */
rte_pktmbuf_free(buf[0]);
```

执行命令：

```
cd dpdk
sudo meson -D examples=all build
cd build
sudo ninja install
cd examples
sudo ./dpdk-skeleton
```

通过 wireshark，我们可以监听来自虚拟机的 UDP 包，且 UDP 包可以正常解析，内容正确。



References

- [1] 深入浅出 dpdk chapter1.7 实例讲解
- [2] 深入浅出 dpdk chapter6.6 Mbuf 与 Mempool
- [3] <https://blog.csdn.net/XuVowkin/article/details/117064512>
- [4] DPDK 总结之常用 API [https://blog.csdn.net/gerald\\_jones/article/details/106600175](https://blog.csdn.net/gerald_jones/article/details/106600175)
- [5] DPDK 基础模块之 rte\_mbuf 详解 <https://www.cnblogs.com/ziding/p/4214499.html>
- [6] <https://en.wikipedia.org/wiki/IPv4>